

ALI AKARMA

Trustworthy AI & Agentic Systems Researcher

Madinah, Saudi Arabia | 443059463@stu.iu.edu.sa

www.aliakarma.codes | [GitHub](https://github.com/aliakarma) | [Google Scholar](https://scholar.google.com/citations?user=...) | [LinkedIn](https://www.linkedin.com/in/aliakarma)

RESEARCH PROFILE

Undergraduate researcher specializing in trustworthy agentic AI, safe reinforcement learning, adversarial robustness, and AI assurance for autonomous systems. Research experience spans LLM-agent security, multi-agent decision-making, human oversight, verification and validation, and safety evaluation of autonomous AI.

RESEARCH INTERESTS

Trustworthy and safe agentic AI; AI safety and alignment; adversarial robustness and security of autonomous LLM systems; safety-constrained multi-agent reinforcement learning; AI verification, assurance, and human oversight; trustworthy multimodal AI.

EDUCATION

B.S. Information Technology — Islamic University of Madinah, Saudi Arabia 2023 – May 2027

- GPA: 4.31 / 5.00 — *Merit-based fully funded scholarship (competitive international selection)*
- Relevant coursework: Machine Learning, Algorithms & Data Structures, Probability & Statistics, Database Systems, Software Engineering
- Thesis interest: Trustworthy and verifiable autonomous AI systems under adversarial conditions

RESEARCH EXPERIENCE

Research Intern — AI V&V Lab, KFUPM Jun 2026 -- Present

- Developed and evaluated HiVA, a value-aware hierarchical multi-agent RL system for wildfire suppression, using Cell2Fire and testing robustness across 16 distribution-shift stress conditions, including cross-region California↔Saudi transfer, rotated wind, ignitions, and asset layouts.
- Contributed to an agentic extension of vnvspec, a machine-readable V&V framework, by developing governance controls for prompt security, tool authorization, memory leakage, human oversight, and multi-agent coordination, with typed schemas and automated validation.

Undergraduate AI Researcher — Islamic University of Madinah Jan 2025 -- Present

- Designed safety-constrained multi-agent reinforcement learning architectures using C-Dec-POMDP and Lagrangian CTDE, achieving 81.5 reward with a 2.3% safety-violation rate across six baselines.
- Analyzed adversarial prompt-injection and jailbreak failure modes in autonomous LLM-agent pipelines; developed LangChain defense middleware with three detection modes and four policy strategies.
- Developed a governance-constrained agentic AI framework incorporating human oversight for safety-critical autonomous decision-making; first author, ICETAS 2026.
- Built 20+ reproducible open-source research systems spanning agentic AI, safety-critical decision-making, AI security, healthcare, and smart infrastructure, with automated testing and archived research artifacts.

RESEARCH PROJECTS & FUNDING

King Salman Center for Disability Research (KSCDR) — *Research Group No. KSRG-2026-321*

- Siddiqui, M. S., Syed, T. A., & Akarma, A. (2026). "ADAPT: An Agentic AI Framework for People with Disabilities and Neurodivergence". *Journal of Disability Research*, 5(2), e20260830.

DOI: [10.57197/JDR-2026-0830](https://doi.org/10.57197/JDR-2026-0830) | GitHub: github.com/aliakarma/ADAPT

- Syed, T. A., Siddiqui, M. S., Akarma, A., & Formisano, A. (2026). "FedAgent-Chain: A Secure Federated and Agentic AI Framework for Multilingual Disability-Inclusive Employment in AI Cities". *Smart Cities*, 9(7), 106.

DOI: [10.3390/smartcities9070106](https://doi.org/10.3390/smartcities9070106) | GitHub: github.com/aliakarma/fedagent-chain

SELECTED PUBLICATIONS

Google Scholar: 107 citations | *h-index*: 8 | *i10-index*: 3 (September 2026).

First-Author Papers

- **Akarma, A.**, Syed, T. A., Khan, M., & Ahmad, A. (2026). Privacy Leakage in Federated Learning: Gradient-Based Client Identity Inference and Defenses for Inertial Sensing in Vehicular Edge Networks. **104th VTC Fall Boston** [arXiv:2609.02971](#)
- **Akarma, A.**, Syed, T.A., Jan, S., Muneer, H., Jilani, A.K. "Governance-Constrained Agentic AI: Blockchain-Enforced Human Oversight for Safety-Critical Wildfire Monitoring," **10th ICETAS 2026**. [arXiv:2604.04265](#)
- **Akarma, A.**, Syed, T.A., Jilani, A.K., Jan, S., Muneer, H., Khan, M.A., Yu, C. "Agents for Agents: An Interrogator-Based Secure Framework for Autonomous Internet of Underwater Things," **10th ICETAS 2026**. [arXiv:2604.04262](#)
- **Akarma, A.**, Syed, T.A., Hameed, D., Muneer, H. "Improving Disability Support and Institutional Performance With AI," Book Chapter, **IGI Global**. [preprints202608.1878](#)

Journal Articles

- Syed, T. A., Akarma, A., Kamal, S., Jan, S., Alkhodre, A. B., & Jamal, A. (2026). Governance-aware autonomous retail coordination in artificial intelligence cities using multi-agent reinforcement learning, blockchain accountability, and federated learning. *Frontiers in Artificial Intelligence*. [10.3389/frai.2026.1942376](#) (Co-First Author)
- Syed TA, **Akarma A**, Alatif A, Naqash MT, Alqurashi A. "Agentic AI-enhanced digital twins for Smart City Civil Infrastructure: A Secure, Autonomous and Auditable Management Framework" **PLOS ONE**. [10.1371/journal.pone.0353610](#) (Co-First Author).
- Jan, S., **Akarma, A.**, Syed, T.A., Muhammad, M. A., & Kamal, S. "EAGF: a four-pillar ethical AI governance framework for trustworthy cybersecurity in 5G renewable energy IoT systems". *Scientific Reports* 2026. [10.1038/s41598-026-63383-5](#).
- Syed, T. A., **Akarma, A.**, Naqash, M. T., Hameed, D., Kamal, S., & Formisano, A. "Agentic AI for Climate-Resilient Cities: A PRISMA-Guided Review and Digital Twin Framework," **MDPI Sustainability**. [10.3390/su18178917](#)
- Syed, T. A., Alshahrani, A., **Akarma, A.**, Khan, S., Nauman, M., Lee, I. E., ... & Ullah, A. "FinNutriAgent (FNA): An Agentic AI for Budget and Nutrition Planning," **ETASR**, 2026. [10.48084/etasr.15640](#)

OPEN-SOURCE RESEARCH SOFTWARE

Selected research systems include reproducible experimental pipelines, automated testing, and archived research artifacts.

- **Trustguard** — Constrained multi-agent reinforcement learning framework for trustworthy runtime permission governance using MAPPO and Lagrangian safety constraints. ([Repository](#))
- **Secure Agent Runtime** — Multimodal security framework for autonomous LLM agents, covering prompt injection, provenance, and text, image, audio, and video threats. ([Repository](#))
- **SAS-Shield** — Safety-assurance framework for agentic AI combining conformal testing, formal verification, control-barrier constraints, adversarial evaluation, and cryptographic auditability. ([Repository](#))

Additional research software and reproducibility artifacts: github.com/aliakarma

TECHNICAL SKILLS

AI/ML: PyTorch, Hugging Face Transformers, scikit-learn, NumPy

Agentic AI & LLMs: LLM orchestration, autonomous agents, prompt-injection defense, adversarial robustness

Reinforcement Learning: MAPPO, CTDE, C-Dec-POMDP, Lagrangian constrained optimization

Trustworthy AI: safety evaluation, failure-mode analysis, human oversight, policy-aligned control, auditability

Research Engineering: Python, Git, Linux, Jupyter, pytest, LaTeX, Zenodo

AWARDS & CERTIFICATIONS

- **Merit-Based Fully Funded Scholarship** — Islamic University of Madinah (competitive international selection, 2023)
- **Certificate of Research Excellence** — ICETAS 2026, for research contribution on Agentic AI governance
- **Certificate of Appreciation** — ICBDT 2025, for presented research contributions

LANGUAGES

Languages: Urdu (native), English (professional proficiency), Arabic (professional proficiency)